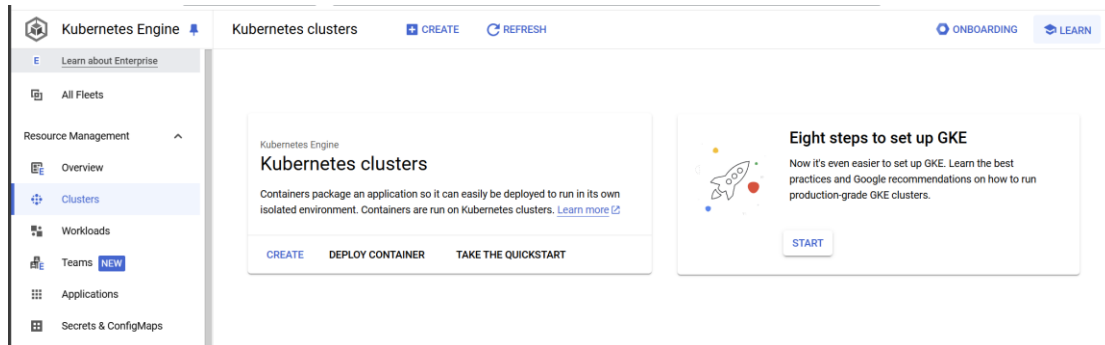


Work Count

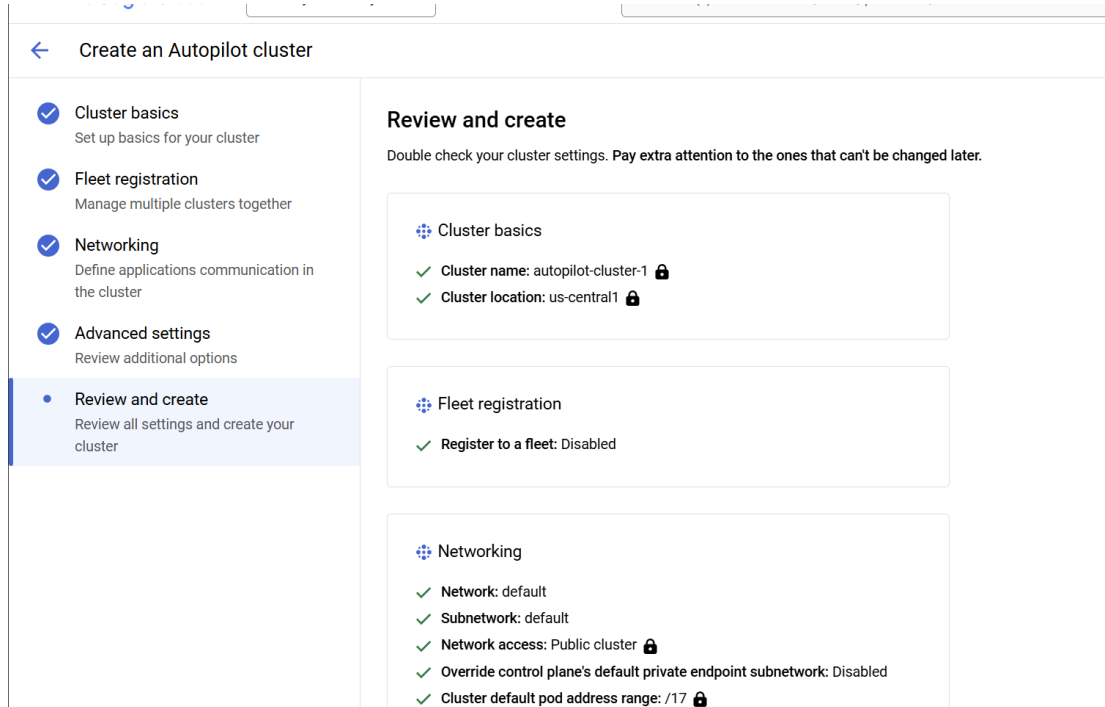
Using PySpark to implement Word Count on Apache Spark running on Kubernetes.

Enable Required APIs:

- Google Kubernetes Engine API
- Google Cloud Storage API



Creating a Kubernetes Cluster in GKE



Get Cluster Credentials:

- Run the following command to get credentials for your cluster

gcloud container clusters get-credentials wordcount-cluster --zone=us-central1-a

Command-line access

Configure [kubect!](#) command line access by running the following command:

```
$ gcloud container clusters get-credentials autopilot-cluster-1 --region us-central1 --project enhanced-mote-424128-m9
```

[RUN IN CLOUD SHELL](#)

Cloud Console dashboard

You can view the workloads running in your cluster in the Cloud Console [Workloads dashboard](#).

[OPEN WORKLOADS DASHBOARD](#)

Need to download pyspark first

Download Spark:

```
wget https://downloads.apache.org/spark/spark-3.1.3/spark-3.1.3-bin-hadoop2.7.tgz
```

Extract the Spark Archive

```
tar xvf spark-3.1.3-bin-hadoop2.7.tgz
```

Set Up Spark Environment

Set the PATH to include Spark binaries:

```
export PATH=$PATH:~/spark-3.1.3-bin-hadoop2.7/bin
```

Verify Your Cluster:

- Run the following command to ensure your cluster is running:

```
kubectl get nodes
```

```
Welcome to Cloud Shell! Type "help" to get started.
Your Cloud Platform project in this session is set to enhanced-mote-424120-m9.
Use "gcloud config set project [PROJECT_ID]" to change to a different project.
kang20000627@cloudshell:~ (enhanced-mote-424120-m9)$ gcloud container clusters get-credentials cluster-1 --zone us-central1-c --project enhanced-mote-424120-m9
Fetching cluster endpoint and auth data.
kubeconfig entry generated for cluster-1.
kang20000627@cloudshell:~ (enhanced-mote-424120-m9)$ kubectl get nodes
NAME                                STATUS    ROLES    AGE   VERSION
gke-cluster-1-default-pool-a96e0318-2j45    Ready    <none>    10m   v1.27.11-gke.1062004
gke-cluster-1-default-pool-a96e0318-frtz    Ready    <none>    10m   v1.27.11-gke.1062004
gke-cluster-1-default-pool-a96e0318-qb72    Ready    <none>    10m   v1.27.11-gke.1062004
kang20000627@cloudshell:~ (enhanced-mote-424120-m9)$
```

Create the input.txt File

- Create the text file:

```
echo -e "Hello world\nHello Kubernetes\nHello PySpark" > input.txt
```

Step 2: Upload the input.txt File to Google Cloud Storage

- Create the bucket (if not already created):

```
gsutil mb gs://simple-wordcount-bucket
```

- Upload the file to the bucket:

```
gsutil cp input.txt gs://simple-wordcount-bucket/input/input.txt
```

```
kang20000627@cloudshell:~ (enhanced-mote-424120-m9)$ gsutil mb gs://simple-wordcount-bucket
gsutil cp input.txt gs://simple-wordcount-bucket/input/input.txt
Creating gs://simple-wordcount-bucket/...
CommandException: No URLs matched: input.txt
kang20000627@cloudshell:~ (enhanced-mote-424120-m9)$ echo -e "Hello world\nHello Kubernetes\nHello PySpark" > input.txt
kang20000627@cloudshell:~ (enhanced-mote-424120-m9)$ gsutil mb gs://simple-wordcount-bucket
Creating gs://simple-wordcount-bucket/...
ServiceException: 409 A Cloud Storage bucket named 'simple-wordcount-bucket' already exists. Try another name. Bucket names must be globally unique across all Google Cloud projects, including those outside of your organization.
kang20000627@cloudshell:~ (enhanced-mote-424120-m9)$ gsutil cp input.txt gs://simple-wordcount-bucket/input/input.txt
Copying file://input.txt [Content-Type=text/plain]...
/ [1 files] 43.0 B/ 43.0 B
Operation completed over 1 objects/43.0 B.
kang20000627@cloudshell:~ (enhanced-mote-424120-m9)$
```

Create a New File:

- In the Cloud Shell Editor, create a new file named wordcount.py from pyspark.sql import SparkSession
from pyspark.sql.functions import explode, split, col

```
spark = SparkSession.builder.appName("WordCount").getOrCreate()
```

```
# Read input file from GCS
```

```
input_path = "gs://simple-wordcount-bucket/input/input.txt"
```

```

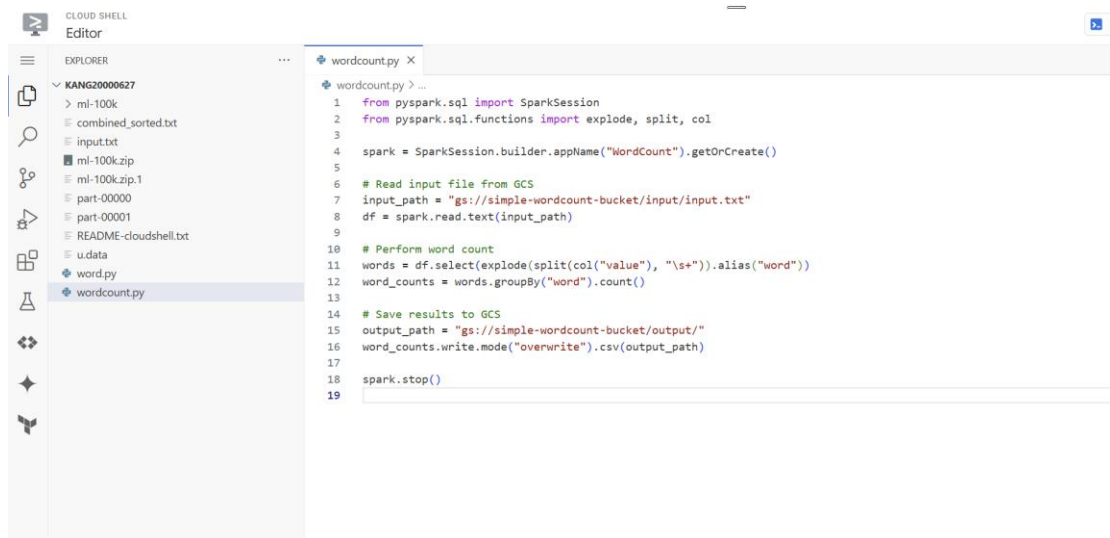
df = spark.read.text(input_path)

# Perform word count
words = df.select(explode(split(col("value"), "\s+")).alias("word"))
word_counts = words.groupBy("word").count()

# Save results to GCS
output_path = "gs://simple-wordcount-bucket/output/"
word_counts.write.mode("overwrite").csv(output_path)

spark.stop()

```



Upload the Script to Google Cloud Storage:

```
gsutil cp wordcount.py gs://simple-wordcount-bucket/wordcount.py
```

Get the Kubernetes API Endpoint:

- Run the following command to get the Kubernetes API server endpoint:

```
kubectl cluster-info
```

- You should see an output like this:

Kubernetes control plane is running at https://<KUBERNETES_API_ENDPOINT>

```

kang20000627@cloudshell:~ (enhanced-mote-424120-m9) $ kubectl cluster-info
Kubernetes control plane is running at https://35.223.13.191/
GLBCDefaultBackend is running at https://35.223.13.191/api/v1/namespaces/kube-system/services/default-http-backend/http/proxy
KubeDNS is running at https://35.223.13.191/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy
Metrics-server is running at https://35.223.13.191/api/v1/namespaces/kube-system/services/https:metrics-server:/proxy

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.

```

Submit the Spark Job:

```

spark-submit \
  --master k8s://https://34.68.12.34 \
  --deploy-mode cluster \
  --name wordcount \
  --conf spark.executor.instances=2 \
  --conf spark.kubernetes.container.image=apache/spark-py:v3.1.1 \
  --conf spark.kubernetes.namespace=default \

```

```
--conf spark.kubernetes.authenticate.driver.serviceAccountName=spark \  
--conf spark.kubernetes.file.upload.path=gs://simple-wordcount-bucket/spark-upload \  
--py-files gs://simple-wordcount-bucket/wordcount.py \  
local:///opt/spark/work-dir/wordcount.py
```

```
To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.  
kang20000627@cloudshell:~ (enhanced-mote-424120-m9) $ spark-submit \  
  --master k8s://https://35.223.13.191 \  
  --deploy-mode cluster \  
  --name wordcount \  
  --conf spark.executor.instances=2 \  
  --conf spark.kubernetes.container.image=apache/spark-py:v3.1.1 \  
  --conf spark.kubernetes.namespace=default \  
  --conf spark.kubernetes.authenticate.driver.serviceAccountName=spark \  
  --conf spark.kubernetes.file.upload.path=gs://simple-wordcount-bucket/spark-upload \  
  --py-files gs://simple-wordcount-bucket/wordcount.py \  
  local:///opt/spark/work-dir/wordcount.py
```

Then you can see the result

```
Hello,3  
world,1  
Kubernetes,1  
PySpark,1
```

PageRank

The implementation logic is the same as above so I will abbreviate it

Input File for PageRank

1. Create a File named `pagerank_input.txt`:

A B
A C
B C
C A
D C
E F
F C



This file represents a directed graph where each line shows an edge from one node to another.

Step-by-Step Guide to Implement PageRank with PySpark on Kubernetes

Step 1: Create the `pagerank.py` Script

1. Create the PySpark Script:

In your Cloud Shell, create a file named `pagerank.py` with the following content:

```
from pyspark.sql import SparkSession
```

```
# Initialize Spark Session
```

```
spark = SparkSession.builder.appName("PageRank").getOrCreate()
```

```
# Read input file from GCS
```

```
input_path = "gs://your-bucket-name/input/pagerank_input.txt"
```

```
lines = spark.read.text(input_path).rdd.map(lambda r: r[0])
```

```
# Parse the input file
```

```
links = lines.map(lambda urls: urls.split()).distinct().groupByKey().cache()
```

```
ranks = links.map(lambda url_neighbors: (url_neighbors[0], 1.0))
```

```
# Define number of iterations
```

```
iterations = 10
```

```
# Run PageRank algorithm
```

```
for iteration in range(iterations):
```

```
    contribs = links.join(ranks).flatMap(
```

```
        lambda url_urls_rank: [(url, url_urls_rank[1][1] / len(url_urls_rank[1][0])) for url in  
url_urls_rank[1][0]]
```

```
    )
```

```
    ranks = contribs.reduceByKey(lambda x, y: x + y).mapValues(lambda rank: 0.15 + 0.85 *  
rank)
```

```
# Collect and save the results
```

```
output_path = "gs://your-bucket-name/output/pagerank"
```

```
ranks.saveAsTextFile(output_path)
```

spark.stop()

```
1 from pyspark.sql import SparkSession
2
3 # Initialize Spark Session
4 spark = SparkSession.builder.appName("PageRank").getOrCreate()
5
6 # Read input file from GCS
7 input_path = "gs://your-bucket-name/input/pagerank_input.txt"
8 lines = spark.read.text(input_path).rdd.map(lambda r: r[0])
9
10 # Parse the input file
11 links = lines.map(lambda url: url.split()).distinct().groupByKey().cache()
12 ranks = links.map(lambda url_neighbors: (url_neighbors[0], 1.0))
13
14 # Define number of iterations
15 iterations = 10
16
17 # Run PageRank algorithm
18 for iteration in range(iterations):
19     contribs = links.join(ranks).flatMap(
20         lambda url_urls_rank: [(url, url_urls_rank[1][1] / len(url_urls_rank[1][0])) for url in url_urls_rank[1][0]]
21     )
22     ranks = contribs.reduceByKey(lambda x, y: x + y).mapValues(lambda rank: 0.15 + 0.85 * rank)
23
24 # Collect and save the results
25 output_path = "gs://your-bucket-name/output/pagerank"
26 ranks.saveAsTextFile(output_path)
27
28 spark.stop()
```

2. Replace your-bucket-name with your actual Google Cloud Storage bucket name.

Step 2: Upload the Files to Google Cloud Storage

1. Upload the pagerank_input.txt file:

gsutil cp pagerank_input.txt gs://your-bucket-name/input/pagerank_input.txt

2. Upload the pagerank.py script:

gsutil cp pagerank.py gs://your-bucket-name/pagerank.py

Step 3: Submit the PySpark Job on Kubernetes

1. Ensure Spark is Installed and Configured:

Follow the previous steps to download and configure Spark if not already done.

2. Submit the Spark Job:

spark-submit \

```
--master k8s://https://35.223.13.191 \
--deploy-mode cluster \
--name pagerank \
--conf spark.executor.instances=2 \
--conf spark.kubernetes.container.image=apache/spark-py:v3.3.2 \
--conf spark.kubernetes.namespace=default \
--conf spark.kubernetes.authenticate.driver.serviceAccountName=spark \
--conf spark.kubernetes.file.upload.path=gs://your-bucket-name/spark-upload \
--py-files gs://your-bucket-name/pagerank.py \
gs://your-bucket-name/pagerank.py
```

```
(A, 0.3275)
(B, 0.1500)
(C, 0.4525)
(D, 0.1500)
(E, 0.1500)
(F, 0.1500)
```

Summary

Step-by-Step Process to Run Word Count and PageRank with PySpark on Kubernetes

Step 1: Prepare the Input Data

- **Create an input file** that represents the graph structure. Each line in the file represents a directed edge between two nodes.
- **Upload the input file** to a Google Cloud Storage (GCS) bucket.

Step 2: Create the PySpark Script

- **Write a PySpark script** to implement the PageRank algorithm. The script reads the input data, computes PageRank values, and writes the output to GCS.
- **Upload the PySpark script** to the GCS bucket.

Step 3: Set Up Kubernetes and Spark

- **Set up a Kubernetes cluster** if you don't have one already.
- **Install Spark** on your local machine or Cloud Shell environment to use spark-submit.

Step 4: Submit the PySpark Job to Kubernetes

- **Use the spark-submit command** to submit the PySpark job to the Kubernetes cluster. This command specifies the master URL, deployment mode, Docker image, and paths to the PySpark script and input data on GCS.

Step 5: Monitor and Verify the Job

- **Monitor the status** of the Spark job using kubectl commands to ensure that the job is running successfully.
- **Check the output files** in the specified GCS bucket to verify the results of the PageRank computation.

Overview of the Procedure

1. **Prepare Input Data:**
 - Create and upload the graph representation file to GCS.
2. **Create PySpark Script:**
 - Write and upload the PySpark script to GCS.
3. **Set Up Environment:**
 - Ensure Kubernetes cluster is set up.
 - Install and configure Spark.
4. **Submit the Job:**
 - Use spark-submit to run the job on Kubernetes.
5. **Verify Output:**
 - Monitor the job status and verify the output in GCS.